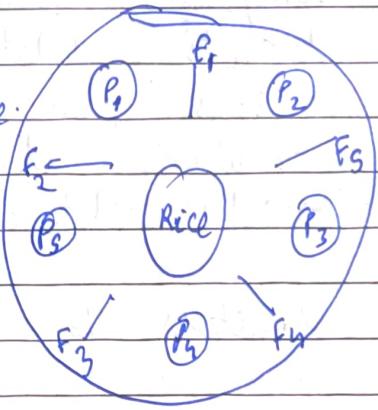


Unit
8.3

Classical Problems of Synchronization:

(1) Dining Philosophers Problem :

- There are 5 philosophers sitting around table with bowl of rice.
- There are 5 chopsticks, one b/w each pair of philosophers.
- To eat philosophers need 2 chopsticks - one on their left & right.
- Each philosopher alternates between think() & eat() operation.
- The challenge is to coordinate their actions so that no two adjacent philosophers try to pick up the same chopsticks at same time causing a deadlock.



Solution :

```

#define N 5
#define LEFT (i+N-1)%N // number of i's left neighbor
#define RIGHT (i+1)%N
#define THINKING 0
#define HUNGRY 1 // Philosopher trying to get fork
#define EATING 2

typedef int Semaphore;
int state[N];
Semaphore mutex = 1;
Semaphore s[N];
    
```

P ₁	S ₁	S ₂
P ₂	S ₂	S ₃
P ₃	S ₃	S ₄
P ₄	S ₄	S ₅
P ₅	S ₅	S ₁

void Philosopher (int i)

{
while (TRUE) {

 think();

 take_forks(i); // Acquire 2 forks or block

 eat();

 Put_forks(i); // Put both forks back to table

}

}

void take_fork (int i)

{

 down(& mutex);

// Enter critical region

 state[i] = HUNGRY;

 test(i);

// try to acquire 2 forks

 up(& mutex);

// exit critical region

 down(& s[i]);

// block if forks were not acquired

}

void Put_forks (int i)

{

 down(& mutex); // enter critical region

 state[i] = HUNGRY; THINKING;

 test(LEFT);

 test(RIGHT);

 up(& mutex); // exit critical region

void test(i) {

{
if (state[i] == HUNGRY && state[LEFT] != EATING && state[RIGHT] !=

EATING)

 state[i] = EATING;

 up(& s[i]);

}

Some database R-W \rightarrow Problem
 W-R \rightarrow Problem
 W-W \rightarrow Problem
 R-R \rightarrow No Problem.

Page No.	
Date	

(2) Readers Writers Problem:

- There is shared data resource that can read concurrently but must be accessed exclusive when writing.
- There are multiple readers that only read data.
- There are multiple writer processes that can update/ modify the data.
- The challenge is to coordinate the readers & writers to allow maximum parallelism but prevent data corruptions.
- Readers increment a read count when starting, decrement on end.
- Writers acquire a write lock, blocking if any readers are active.

Pseudocode (Solution)

```
typedef int semaphore;
Semaphore mutex = 1;
Semaphore db = 1;
int rc = 0;
```

```
{ void reader (void)
```

```
    while (TRUE) {
        down (& mutex); // get exclusive access
        rc = rc + 1; // one more reader now
        if (rc == 1) down (& db);
        up (& mutex); // release exclusive access
        read_database(); // access the data
        down (& mutex);
        rc = rc - 1; // one reader less
        if (rc == 0) up (& db);
    }
}
```

Fair Cases :-

- i) Read First - Writers → Problem
- ii) Write First - Readers → Problem
- iii) Writer - Writers → Problem
- iv) Read - Read → NO Problem

up (& mutex);

use - data - read();

// non critical region
(exit)

}

Void Writer (void)

{

while (TRUE){

think - up - data ();

// non - critical region

down (& db);

// get exclusive access

write - up - data base ();

// update the data

up (& db);

// release

}

}

(3) Producers - Consumers Problems

- There are multiple producer threads / processes generating data items.
- There are multiple consumer threads / processes generating consuming data items.
- The producer & consumer share a fixed size buffer used to store and pass data items.
- Challenge is to synchronize the producers & consumers to ensure:
 - i) Producers cannot add items to a full buffer.
 - ii) Consumers cannot remove items from an empty buffer.
- Counter tracks available slots & items in the buffer.

int count = 0;

{ void producer (void)

{ int itemp;

while (TRUE)

{

 Produce_item (item p);

 while (count == n);

 Buffer [in] = item p;

 in = (in + 1) mod n;

 count = count + 1 →

}

3

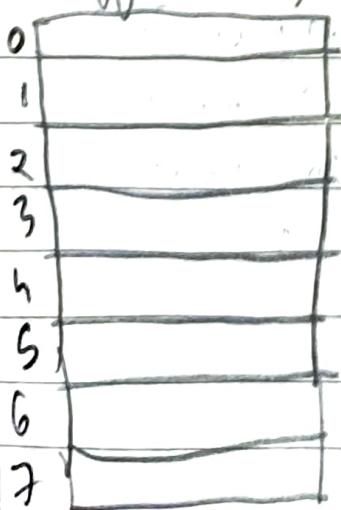
Load R_p, m [count]

INC R_p,

Store m [count], R_p

n = 8

Buffer [0, ..., n - 1]



Count



{ void Consumer(void)

int item C;

while (TRUE)

while (count == 0);

item C = Buffer (Out);

Out = (Out + 1) mod n;

Count = Count - 1;

Process item (item);

}

Load R_C m[Count]

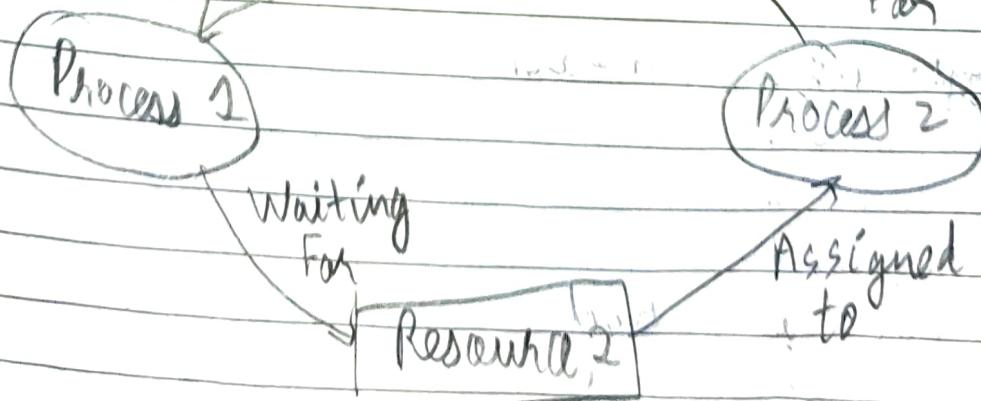
DEC R_C R_C

Store m[Count], R_C

}

Deadlock:

Unit 3, 4
Deadlock refers to a situation where two or more processes are blocked waiting for each other to release of resources, resulting in none of them making progress.



Deadlock Vs Starvation

Deadlock

- It is a process situation where no process gets blocked & no process proceeds.
- Deadlock is an infinite waiting.
- The requested resource is blocked by the other process.
- Every deadlock is always starving.

Starvation

- Starvation is a situation where the low priority process got blocked and high priority processes proceed.
- Starvation is a long waiting but not infinite.
- The requested resource is continuously be used by the higher priority processes.
- Every starvation need not be deadlock.

Principles of Deadlock:

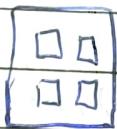
- (i) Mutual Exclusion: Each resource must be either currently assigned to exactly one process or available.
- (ii) Hold and Wait: A process must be holding at least one resource and waiting to acquire additional resources held by other processes.
- (iii) No-Preemption: Resources cannot be forcibly taken away from process they must be released voluntarily by the process holding them.
- (iv) Circular Wait: A set of waiting process must exist such that each process in the set is waiting for a resource held by next process in set; creating circular chain of processes waiting for resources.

V) Resource Allocation Graphs:

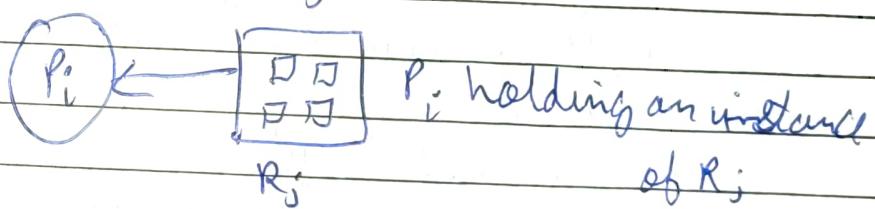
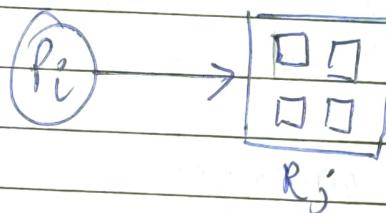
A graph model known as RAG can be used to represent the relationships b/w processes & resources. It helps in visualizing the potential for deadlocks.

Process P_i

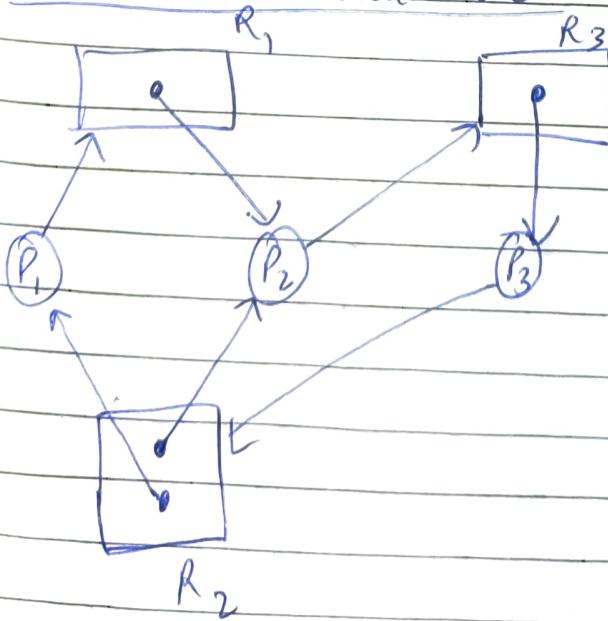
Resource type
(4 instances)



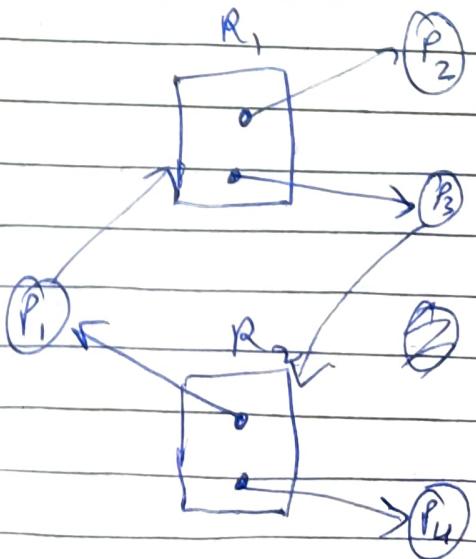
P_i requests instance of R_j



RAG with a Deadlock



RAG with a cycle but no deadlock



Basic Facts

- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycles
 - If only one instance per resource type \Rightarrow deadlock
 - If several " " may or may not

Deadlock Prevention

4 Ways to Prevent Deadlock are:-

- (i) Mutual Exclusion: not required for shareable resources (e.g.: read only files); must hold for non-shareable resources
- (ii) Hold & Wait: must guarantee that whenever a process requests a resource it does not hold any other resources.
 - The process should be required with resources allocated before it begins execution.
 - Low resource utilization starvation possible.
- (iii) No-Preemption:
 - If a process that is holding some resources requests another resource that can not be held immediately allocated to it, then all resources currently being held are released.
- (iv) Circular Wait: To avoid circular wait, we can assign a priority number to each of resource. A process can't request for less priority resource.

◆ Deadlock Avoidance

- The most useful model requires each model process declare the maximum number of resources of each type they need.
- The deadlock avoidance algo. dynamically examines the resource allocation state to ensure that there can be never a circular wait condition.
 - ↳ The resource allocation state of a system can be defined by instances of available & allocated resources, and the maximum instance of the resources demanded by process.

Safe State :

When a process requests an available resource, system must decide if immediate allocation leaves the system in a Safe State.

Basic facts -

If a system is in safe state \rightarrow No deadlock

ii

"unsafe state \Rightarrow possibility of deadlock."

Avoidance \rightarrow Ensure that a system will never enter an unsafe state.

Avoidance Algorithms

Single instance of a resource type - Resource Allocation

Multiple

" - Banker's Algo.

(Deadlock Avoidance)

* Banker's Algorithm (Deadlock Detection)

Total $N = 10$, $B = 5$, $C = 7$

$A \rightarrow CPU$, $B \rightarrow$ Memory (\rightarrow Printers)

(Deadlock avoidance)

(More need - Allocation)

A, B, C
↓
Resource types

Process	Allocation	Max needed	Available			Remaining Need		
			A	B	C	A	B	C
P_1	0 1 0	7 5 3	3 3 2	7	4 3	P_2		
P_2	2 0 0	3 2 2	5 3 2	1	2 2	P_3		
P_3	3 0 2	9 0 2	7 4 3	6	0 0	P_4		
P_4	2 1 1	4 2 2	7 4 5	2	1 1	P_5		
P_5	0 0 2	5 3 3	7 5 5	5	3 1			
	7 2 5		10 5 5					

Available = Total - Allocation.

Safe Sequence
($P_2 \rightarrow P_4 \rightarrow P_5 \rightarrow P_1 \rightarrow P_3$)

No deadlock will occur in this

$$\begin{aligned} \text{Current Availability} &= 10 - 7 = 3 \\ \text{Availability} &= 2 - 2 = 0 \\ &= 5 - 2 = 3 \\ &= 7 - 5 = 2 \end{aligned}$$

All Remaining need \leq Available (Process executes)

A	B	C	
Now	Current Availability		
P_2 Executes & Releases	$\frac{3}{5}$	$\frac{3}{3}$	$\frac{2}{2}$

Banker's Algorithm is static.

$$A \rightarrow 12 \quad D \rightarrow 10$$

$$B \rightarrow 12$$

$$C \rightarrow 8$$

Page No.	
Date	

* Banker's (Resource Request Algorithm)

Process	Allocation				Max				Available			
	A	B	C	D	A	B	C	D	A	B	C	D
P ₀	2	0	0	1	7	2	1	2	3	3	2	1
P ₁	3	1	2	1	5	2	5	2	5	3	2	2
P ₂	2	1	0	3	2	3	1	6	6	6	3	4
P ₃	1	3	1	2	1	4	2	4	7	10	6	6
P ₄	1	4	3	2	3	6	6	5	10	12	8	7
									12	12	8	10

going
in

Need Matrix:-

sequence

Process	Remaining Need			
	A	B	C	D
P ₀	2	2	1	1
P ₁	2	1	3	1
P ₂	0	2	1	3
P ₃	0	1	1	2
P ₄	2	2	3	3

Safe Sequence: P₀ → P₃ → P₄ → P₁ → P₂

- Q. if request from P₂ arrives for (1, 1, 0, 0) can request be immediately granted?

Ans We will update the Allocation vector for P₂ process

& form

$$\begin{bmatrix} 3 & 1 & 2 & 1 \end{bmatrix} \rightarrow \begin{bmatrix} 4 & 2 & 2 & 1 \end{bmatrix} \&$$

Available vector.

$$\begin{bmatrix} 3 & 3 & 2 & 1 \end{bmatrix} \rightarrow \begin{bmatrix} 2 & 2 & 2 & 1 \end{bmatrix} \&$$

Need vector

$$[2 \ 1 \ 3 \ 1] \rightarrow [1 \ 0 \ 3 \ 1].$$

Safe sequence after

this request is = $P_0 P_3 P_4 P_1 P_2$

- Banker's Algo:

- i) Multiple instances
- ii) Each process must a priori claim maximum use.
- iii) When a process requests a resources it must return them in a finite may have to wait.
- iv) When a process gets all its resources it must return them in a finite amount of time.

Data Structure for Banker's Algo :-

let $n = \text{no of processes}$

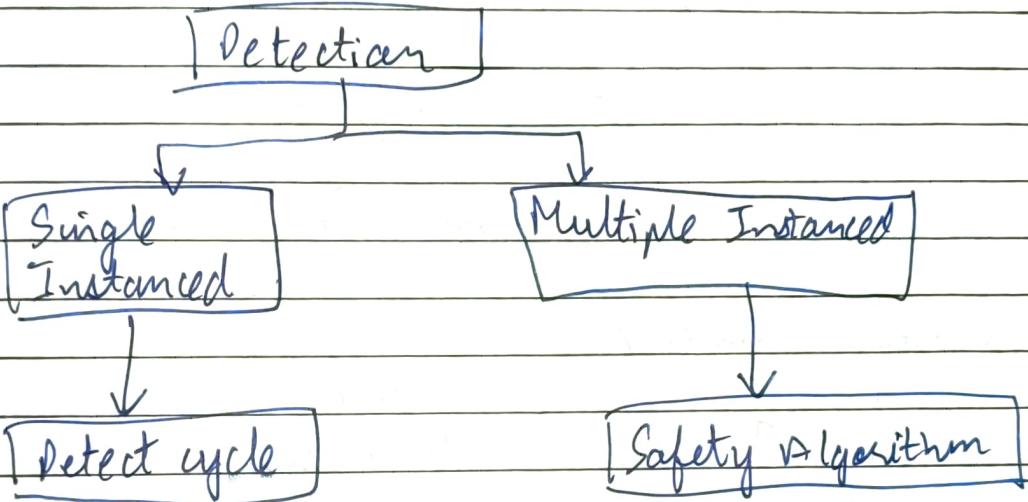
$m = \text{no of resource types}$

- Available : Vector of length m . If available $[j] = k$ there are k instances of resource type R_j available.
- Max : $n \times m$ matrix
- Allocation : $n \times m$ matrix, If Allocation $[i, j] = k$ then P_i is currently allocated k instance
- Need : $n \times m$ matrix. $\text{Need}[i, j] = k$

$$\text{Need}[i, j] = \text{Max}[i, j] - \text{Allocation}[i, j]$$

Deadlock Detection

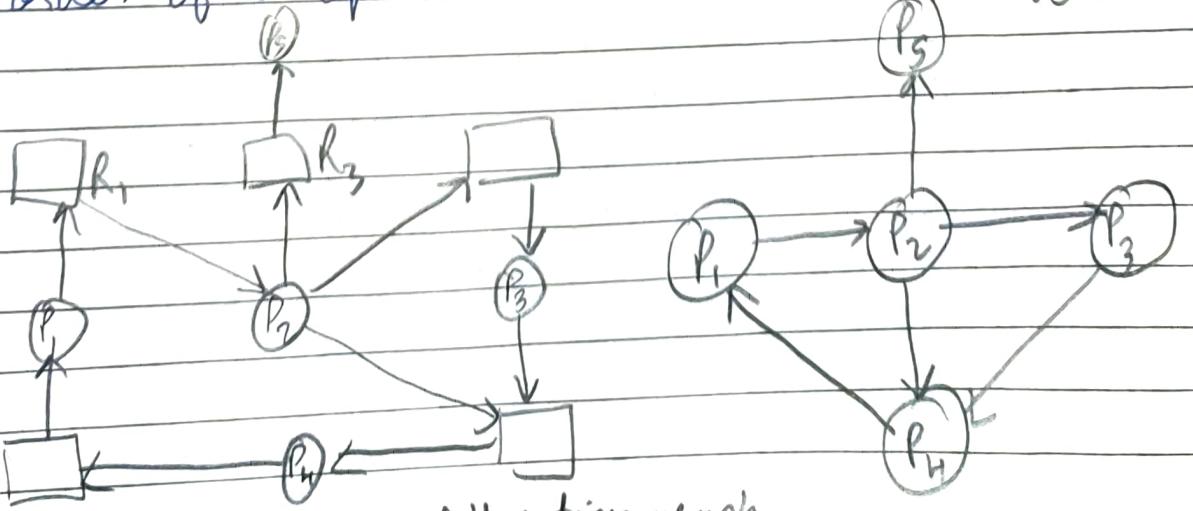
- Deadlock detection involves determining if a deadlock has occurred by analysing state of the system.
- It allows deadlock to occur, then detects it and initiates recovery.



If there is cycle there is a deadlock.

* Single Instance: Maintain wait-for graph
Nodes are processes
 $P_i \rightarrow P_j$ if P_i is waiting for P_j .

An algo to detect a cycle in a graph requires an order of n^2 operations.



(i) Available: A vector of length m

(ii) Allocation: $n \times m$ matrix

(iii) Request: An $n \times m$ matrix

Request $[i][j] = k$, the process P_i is requesting k resources

* Detection Algorithm : (safety algo)

(i) Let work (m) & Finish (n) be vectors.

(a) Work = Available

(b) For $i = 1, 2, \dots, n$ if Allocation $[i] \neq 0$ then
Finish $[i] = \text{False}$, otherwise Finish $[i] = \text{true}$.

(ii) Find an index i such that both :

(a) Finish $[i] == \text{false}$

(b) Request $[i] \leq \text{Work}$

(iii) Work = Work + Allocation $[i]$

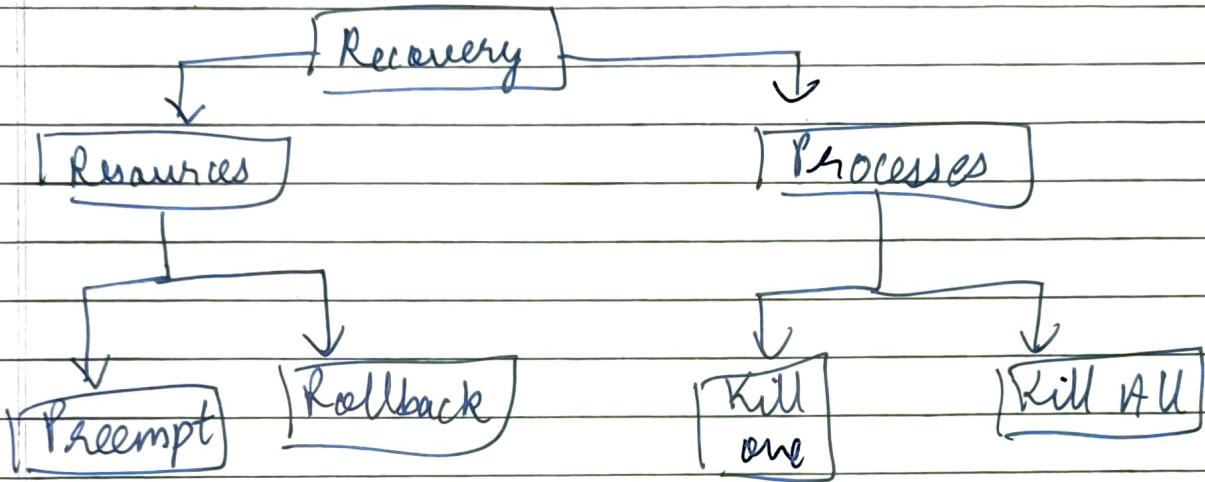
Finish $[i] = \text{true}$

(iv) If Finish $[i] == \text{false}$ for some i , then the system is in deadlocked state.

Algorithm complexity : $O(m \times n^2)$ to detect whether system is in deadlock state.

Deadlock Recovery:

Abort all deadlocked processes.



Recovery :

Preempt - We can snatch one resource from the owner (process) and give it to another process with expectation that it will complete the execution and will release sooner.

Roll back - System passes through various states to get into deadlock state. The OS can roll back the system to previous safe state.

Processes :

Kill a process - The OS kills a process which has done least amount of work until now.

Kill all process - This is not suggestible approach but can be implemented if the problem is serious.