

OPERATING SYSTEMS

MODULE 4

4.1

Memory Management Concepts:

Memory management is a critical aspect of computer systems, involving the organization and control of computer memory. The primary goals include efficient allocation and deallocation of memory, as well as protection to prevent unauthorized access.

Memory Management Requirements:

1. **Allocation:** Assigning portions of memory to different processes.
2. **Deallocation:** Reclaiming memory when it is no longer needed.
3. **Protection:** Ensuring that one process cannot access the memory space of another.

Memory Partitioning:

1. Fixed Partitioning:

- Divides memory into fixed-sized partitions.
- Each partition is assigned to a process.
- Simplicity but can lead to inefficient use of memory, especially if processes have varying memory requirements.

2. Dynamic Partitioning:

- Allocates memory based on the actual needs of processes.
- More flexible than fixed partitioning.
- Requires efficient algorithms to manage variable-sized partitions.

3. Buddy Systems:

- Allocates memory in blocks that are powers of 2.
- If a process needs a certain size, it is allocated the nearest available block size.
- Reduces external fragmentation.

4. Fragmentation:

- External Fragmentation: Free memory scattered in small, non-contiguous blocks.
- Internal Fragmentation: Unused memory within a partition.
- Both types can lead to inefficient memory use and reduced system performance.

5. Paging:

- Divides physical memory into fixed-sized blocks (pages).
- Logical memory is also divided into fixed-sized blocks.
- Helps reduce external fragmentation.
- Page tables are used for mapping logical addresses to physical addresses.

6. Segmentation:

- Divides logical memory into variable-sized segments.
- Each segment corresponds to a different type of data.
- Simplifies memory management but may lead to internal fragmentation.

7. Address Translation:

- Involves mapping logical addresses to physical addresses.
- Performed by hardware or operating system.
- Ensures proper memory access for processes.

In summary, memory management strategies aim to optimize resource utilization and ensure efficient, secure operation of computer systems. Each concept addresses specific challenges and trade-offs in achieving these objectives.

4.2

Placement Strategies: First Fit, Best Fit, Next Fit & Worst Fit

1. First Fit:

- Algorithm: The first available block that is large enough to accommodate the process is allocated.
- Advantages: Simple and easy to implement.
- Disadvantages: May lead to fragmentation, as the first available block might not be the best fit for the process.

2. Best Fit:

- Algorithm: Allocates the smallest available block that is large enough for the process.
- Advantages: Tends to minimize fragmentation by using the smallest block possible.
- Disadvantages: Can be slower than First Fit since it searches for the best match.

3. Next Fit:

- Algorithm: Similar to First Fit but starts searching for a suitable block from the last allocated position and moves forward.
- Advantages: Simplifies bookkeeping, as it continues the search from the previous allocation.
- Disadvantages: Similar to First Fit, it may result in fragmentation.

Comparison:

- Efficiency: First Fit and Next Fit are faster as they find the first available block, while Best Fit might involve more searching.
- Fragmentation: Best Fit is designed to minimize fragmentation,

Worst Fit:

Algorithm:

- Allocates the largest available block to the process.
- The idea is to leave the largest possible free block for future allocations.

Advantages:

- Can help in reducing long-term fragmentation by leaving larger free spaces.
- Similar to Best Fit in terms of trying to optimize memory utilization.

Disadvantages:

- May lead to increased fragmentation, especially if smaller blocks need to be allocated in the future.
- Requires searching for the largest available block, which can be computationally expensive.

Comparison with Other Placement Strategies:

- Efficiency: Worst Fit tends to be slower than First Fit and Next Fit due to the need to find the largest available block.
- Fragmentation: While it aims to leave larger free spaces, it can result in increased fragmentation over time, especially if processes of varying sizes are involved.
- Implementation Complexity: Similar to Best Fit, it is more complex than First Fit and Next Fit due to the searching involved.

In summary, Worst Fit aims to maximize the available free space by allocating the largest block to a process. However, this approach has trade-offs, and its efficiency and effectiveness depend on the specific characteristics of the processes and memory allocations in a given system.

4.3

Virtual Memory:

1. Cache Memory Organization:

- Purpose: To store frequently accessed data for quicker retrieval.
- Levels: Typically organized into multiple levels (L1, L2, L3) to balance speed and capacity.

2. Cache Architecture (L1, L2, L3):

- L1 Cache: Closest to the CPU, small in size, and fastest.
- L2 Cache: Larger and slightly slower than L1.
- L3 Cache: Shared among multiple cores, larger but slower than L2.

3. Address Mapping Techniques:

- Direct Mapping: Each block of main memory maps to exactly one cache line.
- Set-Associative Mapping: Allows a block to be placed in one of several possible sets.
- Fully Associative Mapping: Any block can be placed in any cache location.

4. Cache Coherence:

- Ensures: Consistency of data in multiple caches that store copies of the same data.
- Protocols: MESI (Modified, Exclusive, Shared, Invalid) is a common protocol.

5. Swapping Issues:

- Swapping: Moving a process between main memory and secondary storage.
- Issues: Can be slow due to disk I/O, and excessive swapping can degrade performance.

6. Thrashing:

- Definition: High paging activity, where the majority of processor time is spent in servicing page faults.
- Causes: Insufficient RAM, high degree of multiprogramming, improper page replacement algorithms.

7. VM with Paging:

- Paging: Dividing physical memory into fixed-sized pages.
- Advantages: Efficient use of memory, reduces external fragmentation.

8. Page Table Structure:

- Translation: Maps virtual addresses to corresponding physical addresses.
- Structure: Contains page table entries with information like frame number and status bits.

9. Inverted Page Table:

- Design: One entry per frame, mapping to the process and page number.
- Advantages: Reduces memory overhead, especially for large address spaces.

10. Translation Lookaside Buffer (TLB):

- Purpose: A cache for page table entries, reducing the time to access the page table.
- Efficiency: Accelerates address translation by storing frequently accessed mappings.

11. Page Size:

- Definition: The amount of data in a single page.
- Impact: Larger page sizes can improve memory access speed but may lead to more internal fragmentation.

12. VM with Segmentation:

- Segmentation: Divides a program into different segments (code, data, stack).
- Advantages: Provides logical structure, simplifies memory management.

13. VM with Combined Paging and Segmentation:

- Integration: Combines the benefits of both paging and segmentation.
- Flexibility: Allows for efficient use of memory and provides a logical structure for processes.

In summary, virtual memory involves complex mechanisms like cache organization, addressing techniques, coherence protocols, and handling issues like swapping and thrashing to optimize overall system performance.

4.4

Page Replacement Policies:

1. FIFO (First-In-First-Out):

- Algorithm: Replaces the oldest page in memory.

- Advantages: Simple to implement.
- Disadvantages: May not consider the usage patterns; old pages may still be relevant.

2. LRU (Least Recently Used):

- Algorithm: Replaces the page that has not been used for the longest time.
- Advantages: Considers usage patterns, tends to retain more relevant pages.
- Disadvantages: Implementation can be more complex, and maintaining accurate usage information can be resource-intensive.

3. Optimal:

- Algorithm: Replaces the page that will not be used for the longest time in the future.
- Advantages: Theoretical optimum for minimizing page faults.
- Disadvantages: Impractical for real systems as it requires knowledge of future page accesses.

4.5

Windows/Linux Memory Management:

Windows:

- Paging: Utilizes a demand-paging system, bringing in pages only when needed.
- Page Replacement: Employs a combination of algorithms, including FIFO and LRU.
- Memory Compression: Windows may compress pages in memory to reduce the need for paging to disk.

Linux:

- Swapping: Similar to Windows, Linux swaps pages to disk when necessary.
- Page Replacement: Linux uses various algorithms, including the Completely Fair Scheduler (CFS) for process scheduling.
- Page Cache: Linux employs a page cache to speed up file access by storing file pages in memory.

Comparison:

- Windows vs. Linux: Both use demand-paging and employ various page replacement policies, but specific algorithms and implementations may differ.
- Commonality: Both aim to optimize memory usage, minimize page faults, and ensure efficient system performance.

In summary, page replacement policies like FIFO, LRU, and Optimal play a crucial role in managing virtual memory. Windows and Linux employ similar concepts but may differ in the specific algorithms and mechanisms they use for memory management.

MODULE 5

5.1 - I/O Management:

I/O Devices:

- Types: Input devices (e.g., keyboard, mouse), output devices (e.g., display, printer), and storage devices (e.g., hard drives, SSDs).
- Characteristics: Varying speeds, data rates, and communication protocols.

Characteristics of Devices:

- Speed: Some devices are slower (e.g., printers) compared to faster ones (e.g., RAM).
- Volatility: Some devices are non-volatile (e.g., hard drives) while others are volatile (e.g., RAM).
- Capacity: Devices differ in their storage capacities.

OS Design Issues for I/O Management:

- Device Independence: OS abstracts device details for application programs.
- Buffering: Use of buffers to enhance I/O performance.
- Error Handling: Strategies for detecting and managing I/O errors.
- Concurrency Control: Handling multiple I/O requests concurrently.

I/O Buffering:

- Purpose: To cope with speed mismatches between I/O devices and CPU.
- Types: Single buffering, double buffering, circular buffering.
- Advantages: Enhances system performance by overlapping I/O and CPU operations.

5.2 - Disk Scheduling:

Disk Scheduling Algorithms:

- **FCFS (First-Come-First-Serve)**: Serves requests in the order they arrive. Simple but may result in poor performance.
- **SSTF (Shortest Seek Time First)**: Prioritizes the request with the shortest seek time. Reduces overall seek time but can lead to starvation.
- **SCAN**: Moves the disk arm from one end to the other, serving requests along the way. Helps reduce starvation but may not be optimal.
- **C-SCAN (Circular SCAN)**: Similar to SCAN, but jumps to the other end after reaching one end. Reduces wait times for requests near the edges.
- **LOOK**: Similar to SCAN but doesn't go all the way to the end, reducing the arm movement.
- **C-LOOK (Circular LOOK)**: Similar to LOOK but jumps to the other end after reaching one end.

5.3 - File Management:

Concepts

- File: A file is a collection of related information that is stored on a secondary storage device (such as a hard drive, solid-state drive, or optical disc).
- File system: A file system is a collection of data structures that organize and manage files on a secondary storage device.
- Directory: A directory is a collection of files and other directories.
- Path: A path is a sequence of directory names that specifies the location of a file or directory.

File Organization and access

- File organization: Files are organized into directories and subdirectories. This allows users to group related files together and easily find them.
- File access: Files can be accessed using a path or a file handle. A path is a sequence of directory names that specifies the location of a file. A file handle is a unique identifier that the operating system assigns to a file when it is opened.

File Directories

- Directory structure: The directory structure is a hierarchical tree of directories and files.
- Directory entries: Each directory entry contains the name of a file or directory and a pointer to its location on the secondary storage device.
- Directory operations: The operating system provides a number of directory operations, such as creating, deleting, and renaming directories.

File Sharing

- File sharing: File sharing allows multiple users to access the same file.
- File access control: The operating system can control who can access a file and what operations they can perform on it.
- File sharing protocols: There are a number of file sharing protocols, such as NFS, SMB, and FTP.

File allocation

- File allocation: File allocation is the process of assigning blocks of disk space to a file.
- File allocation strategies: There are a number of file allocation strategies, such as contiguous allocation, linked allocation, and indexed allocation.
- File allocation tables: File allocation tables are used to keep track of which blocks of disk space are allocated to a file.

Secondary Storage Management

- Secondary storage management: Secondary storage management is the process of managing secondary storage devices, such as hard drives and solid-state drives.
- Free space management: Free space management is the process of keeping track of which blocks of disk space are not in use.
- Disk defragmentation: Disk defragmentation is the process of rearranging files on a disk to improve performance.

Security

- File security: File security is the process of protecting files from unauthorized access, modification, or deletion.
- File access control lists (ACLs): ACLs are used to control who can access a file and what operations they can perform on it.
- Encryption: Encryption can be used to protect the contents of a file from unauthorized access.

5.4 - Windows File System:

FAT

- FAT: File Allocation Table is an early file system that was used in MS-DOS and early versions of Windows.
- Features:
 1. Simple and efficient for small disks
 2. Supports up to 512MB partitions
- Limitations:
 1. Not scalable to large disks
 2. Does not support long file names

FAT32

- FAT32: FAT32 is an updated version of FAT that supports larger disks and longer file names.
- Features:
 1. Supports up to 2TB partitions
 2. Supports long file names
- Limitations:

1. Not as efficient as NTFS
2. Does not support features such as file encryption and disk quotas

NTFS

- NTFS: New Technology File System is a modern file system that is used in Windows 2000 and later.
- Features:
 1. Supports large disks and long file names
 2. Supports features such as file encryption and disk quotas
 3. More efficient than FAT and FAT32
- Limitations:
 1. More complex than FAT and FAT32
 2. Not as compatible with older operating systems

ReFS

- ReFS: Resilient File System is a new file system that was first introduced in Windows Server 2012.
- Features:
 1. Designed for high-availability storage
 2. Supports features such as automatic data integrity checking and repair
 3. More efficient than NTFS for large files

5.5 - Linux File System:

ext-2

- ext-2: Extended Filesystem 2 is an early file system that was used in Linux.
- Features:
 1. Simple and efficient
 2. Supports up to 2TB partitions
- Limitations:
 1. Not scalable to large disks
 2. Does not support features such as file journaling and file system quotas

ext-3

- ext-3: Extended Filesystem 3 is an updated version of ext-2 that supports file journaling.
- Features:
 1. Supports up to 2TB partitions
 2. Supports file journaling
 3. More efficient than ext-2
- Limitations:

1. Not as efficient as ext-4
2. Does not support features such as file system quotas and extended attributes

ext-4

- ext-4: Extended Filesystem 4 is a modern file system that is used in most Linux distributions today.
- Features:
 1. Supports large disks and long file names
 2. Supports features such as file journaling, file system quotas, and extended attributes
 3. More efficient than ext-2 and ext-3
- Limitations:
 1. More complex than ext-2 and ext-3
 2. Not as compatible with older operating systems

reiserFS

- reiserFS: Reiser File System is a file system that is known for its high performance.
- Features:
 1. Designed for high performance
 2. Supports features such as file journaling, file system quotas, and extended attributes
 3. More efficient than ext-2, ext-3, and ext-4
- Limitations:
 1. Not as compatible with other Linux distributions
 2. Does not support features such as large file support and encryption

XFS

- XFS: SGI Extended File System is a file system that is designed for large-scale file systems.
- Features:
 1. Designed for large-scale file systems
 2. Supports features such as file journaling, file system quotas, and extended attributes
 3. More efficient than ext-2, ext-3, and ext-4
- Limitations:
 1. Not as compatible with other Linux distributions
 2. Does not support features such as large file support and encryption

JFS

- JFS: Journaled File System is a file system that is used in AIX and IBM i operating systems.
- Features:
 1. Designed for high performance and reliability
 2. Supports features such as file journaling, file system quotas, and extended attributes

3. More efficient than ext-2, ext-3, and ext-4
 - Limitations:
 1. Not as compatible with other operating systems
 2. Does not support features such as large file support and encryption